

Native Libraries Guide

Table of contents

1 Purpose.....	2
2 Components.....	2
3 Usage.....	2
4 Supported Platforms.....	3
5 Building Native Hadoop Libraries.....	3
5.1 Notes.....	3
6 Loading native libraries through DistributedCache	4

1 Purpose

Hadoop has native implementations of certain components for reasons of both performance and non-availability of Java implementations. These components are available in a single, dynamically-linked, native library. On the *nix platform it is *libhadoop.so*. This document describes the usage and details on how to build the native libraries.

2 Components

Hadoop currently has the following [compression codecs](#) as the native components:

- [zlib](#)
- [gzip](#)
- [bzip2](#)

Of the above, the availability of native hadoop libraries is imperative for the gzip and bzip2 compression codecs to work.

3 Usage

It is fairly simple to use the native hadoop libraries:

- Take a look at the [supported platforms](#).
- Either [download](#) the pre-built 32-bit i386-Linux native hadoop libraries (available as part of hadoop distribution in `lib/native` directory) or [build](#) them yourself.
- Make sure you have any of or all of **>zlib-1.2**, **>gzip-1.2**, and **>bzip2-1.0** packages for your platform installed; depending on your needs.

The `bin/hadoop` script ensures that the native hadoop library is on the library path via the system property `-Djava.library.path=<path>`.

To check everything went alright check the hadoop log files for:

```
DEBUG util.NativeCodeLoader - Trying to load the custom-built native-hadoop library...
```

```
INFO util.NativeCodeLoader - Loaded the native-hadoop library
```

If something goes wrong, then:

```
INFO util.NativeCodeLoader - Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
```

4 Supported Platforms

Hadoop native library is supported only on *nix platforms only. Unfortunately it is known not to work on [Cygwin](#) and [Mac OS X](#) and has mainly been used on the GNU/Linux platform.

It has been tested on the following GNU/Linux distributions:

- [RHEL4/Fedora](#)
- [Ubuntu](#)
- [Gentoo](#)

On all the above platforms a 32/64 bit Hadoop native library will work with a respective 32/64 bit jvm.

5 Building Native Hadoop Libraries

Hadoop native library is written in [ANSI C](#) and built using the GNU autotools-chain (autoconf, autoheader, automake, autoscan, libtool). This means it should be straightforward to build them on any platform with a standards compliant C compiler and the GNU autotools-chain. See [supported platforms](#).

In particular the various packages you would need on the target platform are:

- C compiler (e.g. [GNU C Compiler](#))
- GNU Autotools Chain: [autoconf](#), [automake](#), [libtool](#)
- zlib-development package (stable version >= 1.2.0)

Once you have the pre-requisites use the standard `build.xml` and pass along the `compile.native` flag (set to `true`) to build the native hadoop library:

```
$ ant -Dcompile.native=true <target>
```

The native hadoop library is not built by default since not everyone is interested in building them.

You should see the newly-built native hadoop library in:

```
$ build/native/<platform>/lib
```

where `<platform>` is combination of the system-properties: `${os.name}-${os.arch}-${sun.arch.data.model}`; for e.g. `Linux-i386-32`.

5.1 Notes

- It is **mandatory** to have the `zlib`, `gzip`, and `bzip2` development packages on the target platform for building the native hadoop library; however for deployment it is sufficient to install one of them if you wish to use only one of them.

- It is necessary to have the correct 32/64 libraries of both zlib depending on the 32/64 bit jvm for the target platform for building/deployment of the native hadoop library.

6 Loading native libraries through DistributedCache

User can load native shared libraries through [DistributedCache](#) for *distributing* and *symlinking* the library files

Here is an example, describing how to distribute the library and load it from map/reduce task.

1. First copy the library to the HDFS.
`bin/hadoop fs -copyFromLocal mylib.so.1 /libraries/
mylib.so.1`
2. The job launching program should contain the following:
`DistributedCache.createSymlink(conf);
DistributedCache.addCacheFile("hdfs://host:port/libraries/
mylib.so.1#mylib.so", conf);`
3. The map/reduce task can contain:
`System.loadLibrary("mylib.so");`